


Result throughput per query template



bar = geometric mean over instances of the per-instance median; whisker = slowest to fastest instance. Each panel has its own scale. Hatched / $\geq$ = a lower bound: some instances produced nothing before the timeout and are held at it.